

Daniel Schumacher VisCom HW 1

AUTHOR

HDD249 Dan Schumacher

Homework 1

Set up

```
library(tidyverse)
library(scales)
```

Read `tech_stocks_csv.zip` into your R session using `read_csv()` function from `readr` package. Store the resulting object in `d1`.

Rows: 93 Columns: 981

— Column specification —————

Delimiter: ",",

chr (34): gvkey, indfmt, consol, popsrc, datafmt, tic, cusip, conm, acctch...

dbl (436): fyear, ajex, ajp, currtr, fyr, ismod, ltcm, pddur, scf, src, upd...

lgl (506): adrr, bspr, curuscn, ogm, stalt, udpl, acco, accrt, acoxar, acql...

date (5): datadate, apdedate, fdate, pdate, ipodate

! Use ``spec()`` to retrieve the full column specification for this data.

! Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

Q1

Print a data frame with the medians of `at`, `emp`, and `xrd`.

```
d1 %>%
  select(at, emp, xrd) %>%
  summarize(
    median_at = median(at),
    median_emp = median(emp),
    median_xrd = median(xrd)
  )
```

A tibble: 1 × 3

	median_at	median_emp	median_xrd
	<dbl>	<dbl>	<dbl>
1	93798	72.0	7754

Q2

Print a data frame with the means of `sale`, `oibdp`, and `xrd` for Apple, Meta, and Tesla. For this, you will need to follow these steps:

1. Filter only the observations pertaining to Apple, Meta, and Tesla
2. Group by `conm`
3. Summarize `sale`, `oibdp`, and `xrd` to get their means
4. Output it as a data frame by using `as.data.frame()` function.

```
d1 %>%
  filter(
    conm %in% c('APPLE INC',
                'META PLATFORMS INC',
                'TESLA INC')
  ) %>%
  group_by(conm) %>%
  summarise(
    mean_sale = mean(sale),
    mean_oibdp = mean(oibdp),
    mean_xrd = mean(xrd)
  ) %>%
  as.data.frame()
```

	conm	mean_sale	mean_oibdp	mean_xrd
1	APPLE INC	224763.08	71468.923	11167.000
2	META PLATFORMS INC	43413.00	20110.923	9754.923
3	TESLA INC	18585.48	2672.776	1094.280

```
### same thing but using across
(dlmeans <- d1 %>%
  filter(
    conm %in% c('APPLE INC',
                'META PLATFORMS INC',
                'TESLA INC')
  ) %>%
  group_by(conm) %>%
  summarise(
    across(c(sale,oibdp,xrd), mean)
  ) %>%
  as.data.frame())
```

	conm	sale	oibdp	xrd
1	APPLE INC	224763.08	71468.923	11167.000

```
2 META PLATFORMS INC 43413.00 20110.923 9754.923
3      TESLA INC 18585.48 2672.776 1094.280
```

Q3

Round all the numeric variables in the above data frame to 1 decimal place. Output as a data frame using `as.data.frame()` function.

For rounding, you will have to use `mutate`, `across`, and `where` functions from `dplyr` package

```
d1means %>%
  mutate(
    across(
      where(is.numeric),
      round, 1)
  )
```

Warning: There was 1 warning in `mutate()`.

! In argument: `across(where(is.numeric), round, 1)`.

Caused by warning:

! The `...` argument of `across()` is deprecated as of dplyr 1.1.0.

Supply arguments directly to `.fns` through an anonymous function instead.

```
# Previously
```

```
across(a:b, mean, na.rm = TRUE)
```

```
# Now
```

```
across(a:b, \(x) mean(x, na.rm = TRUE))
```

```
      conm    sale  oibdp    xrd
1      APPLE INC 224763.1 71468.9 11167.0
2 META PLATFORMS INC 43413.0 20110.9 9754.9
3      TESLA INC 18585.5 2672.8 1094.3
```

Q4

Many advertising values are missing. The missing code in R is `NA`. We can get the total number of missing values for advertising quite easily by running the following function: `sum(is.na(d1$xad))`

In the finance literature, a common (but incorrect) practice is to assume that the missing advertising is 0. We will use this adjustment to `xad` and create a new variable `adv` and save it in a new object `d2`.

```
#how many NAs do we have? ----> 13
sum(is.na(d1$xad))
```

```
[1] 13
```

```
#lets replace them with 0s
d2 <- d1 %>%
  #           ifelse(test,   yes-val, no-val)
  mutate(adv = ifelse(is.na(xad), 0, xad))

# we should have 0 NAs now let's check
sum(is.na(d2$adv))
```

[1] 0

Q5

Using `d2`, create the following variables and print first 8 rows for NVidia and the new columns along with `conm` and `datadate`:

1. Return on assets (`roa`) = `oibdp / at`
2. Free cash flow (`fcf`) = `oancf / che`
3. Strategic emphasis (`strat_emph`) = `(adv - xrd) / at`

```
d2 <- d2 %>%
  mutate(
    roa = oibdp/at,
    fcf = oancf/che,
    strat_emph = (adv - xrd) / at
  )

d2 %>%
  select(conm, datadate, roa, fcf, strat_emph) %>%
  filter(conm == 'NVIDIA CORP')
```

A tibble: 14 × 5

	conm	datadate	roa	fcf	strat_emph
	<chr>	<date>	<dbl>	<dbl>	<dbl>
1	NVIDIA CORP	2010-01-31	0.0926	0.282	-0.249
2	NVIDIA CORP	2011-01-31	0.129	0.271	-0.187
3	NVIDIA CORP	2012-01-31	0.158	0.291	-0.179
4	NVIDIA CORP	2013-01-31	0.142	0.221	-0.177
5	NVIDIA CORP	2014-01-31	0.107	0.179	-0.182
6	NVIDIA CORP	2015-01-31	0.136	0.196	-0.187
7	NVIDIA CORP	2016-01-31	0.149	0.233	-0.178
8	NVIDIA CORP	2017-01-31	0.217	0.246	-0.147
9	NVIDIA CORP	2018-01-31	0.303	0.493	-0.158
10	NVIDIA CORP	2019-01-31	0.309	0.504	-0.177
11	NVIDIA CORP	2020-01-31	0.187	0.437	-0.163
12	NVIDIA CORP	2021-01-31	0.205	0.504	-0.136

13	NVIDIA	CORP	2022-01-31	0.254	0.429	-0.119
14	NVIDIA	CORP	2023-01-31	0.174	0.424	-0.178

Q6

You want to know how many profitable years each of the sample company experienced. For this, follow these steps:

1. Create an indicator variable (dummy variable) called `profit_ind` such that when `oibdp > 0` this variable is 1. Otherwise it is 0.

```
d2 <- d2 %>%
  mutate(
    profit_ind = ifelse(oibdp > 0, 1, 0)
  )

#let's check it
d2$profit_ind
```

[illegible]

2. Group by company names
3. Summarize `profit_ind` by taking its sum. Also, get the total number of observations for each company.

```
d2 %>%
  group_by(conm) %>%
  summarise(
    profit_years = sum(profit_ind),
    # I found this answer on S.O.
    # why doesn't ty = count(profit_ind) work?
    total_years = n(), .groups = 'drop')
```

```
# A tibble: 7 × 3
  comn          profit_years total_years
<chr>          <dbl>         <int>
1 ALPHABET INC          13           13
2 AMAZON.COM INC        13           13
3 APPLE INC             13           13
4 META PLATFORMS INC    13           13
5 MICROSOFT CORP        14           14
6 NVIDIA CORP           14           14
7 TESLA INC             9            13
```

Q7

Find the average annual stock returns of all the companies. Follow these steps:

```
d2 %>%
# Arrange the data set by conm and datadate.
  arrange(conm, datadate) %>%

# Group by conm
  group_by(conm) %>%

# Calculate stock return stk_ret by taking the difference between mkvalt (Market value of equity)
  mutate(
    stk_ret =
      (mkvalt - lag(mkvalt, n = 1)) / lag(mkvalt, n = 1)
  ) %>%

# Calculate stock return stk_ret2 by taking the difference between prcc_f (Stock price at the end
  mutate(
    stk_ret2 =
      (prcc_f - lag(prcc_f, n = 1)) / lag(prcc_f, n = 1)
  ) %>%

# Summarize to get the mean of the stock returns stk_ret_mean and stk_ret2_mean.

  summarize(
    stk_ret_mean = percent(mean(stk_ret, na.rm = T), 0.01),
    stk_ret_mean2 = percent(mean(stk_ret2, na.rm = T), 0.01)
  )
```

A tibble: 7 × 3

	conm	stk_ret_mean	stk_ret_mean2
	<chr>	<chr>	<chr>
1	ALPHABET INC	19.87%	11.00%
2	AMAZON.COM INC	29.98%	24.60%
3	APPLE INC	24.31%	5.62%
4	META PLATFORMS INC	29.45%	26.84%
5	MICROSOFT CORP	22.61%	24.01%
6	NVIDIA CORP	57.76%	43.89%
7	TESLA INC	115.37%	43.60%

```
# Display the average stock returns in percentage format.
```

Not graded: The average stock returns calculated using these two measures are very different. Which of these is correct?

1. My guess is `stk_ret` (not `stk_ret2`) because `stk_ret2` is yearly whereas `stk_ret` is on a more regular interval.

Q8

In many statistical and machine learning applications, we use scaled variables instead of the original variables. A scaled variable is typically created by subtracting the sample mean of the variable from the variable and dividing it by its standard deviation. There is a `scale()` function in base R which can directly do it.

You want to create a scaled variable for `sale` but separately for each company. Therefore, you can't use the mean and standard deviation of `sale` for the entire sample. Instead, you have to calculate these statistics for each company separately and then create a scaled variable. Follow these steps:

1. Group by `conm`
2. Summarize `sale` to get the mean (`sale_mean`) and the standard deviation (`sale_sd`)
3. Assign this dataframe to `d2_sum`
4. Join `d2` and `d2_sum` by `conm`
5. Create `sale_scaled` by subtracting `sale_mean` from `sale` and dividing this difference by `sale_sd`

Print the first 10 rows for Tesla with `conm`, `sale`, `sale_scaled`

```
#steps 1-3
(d2_sum <- d2 %>%
  group_by(conm) %>%
  summarize(
    sale_mean = mean(sale, na.rm = T),
    sale_sd = sd(sale, na.rm = T)
  ))
```

```
# A tibble: 7 × 3
  conm                sale_mean sale_sd
  <chr>                <dbl>   <dbl>
1 ALPHABET INC         118540.  82029.
2 AMAZON.COM INC       200842. 164507.
3 APPLE INC            224763.  92452.
4 META PLATFORMS INC   43413   42337.
5 MICROSOFT CORP       114056.  48604.
6 NVIDIA CORP           9914.   8203.
7 TESLA INC            18585.  24698.
```

```
# step 4-6
left_join(d2, d2_sum, by = 'conm') %>%
```

```

mutate(
  sale_scaled = (sale - sale_mean)/sale_sd
) %>%

filter(conm == 'TESLA INC') %>%

select(conm, sale, sale_scaled, sale_mean, sale_sd) %>%

as.data.frame() %>%
head(10)

```

	conm	sale	sale_scaled	sale_mean	sale_sd
1	TESLA INC	116.744	-0.7477853	18585.48	24697.92
2	TESLA INC	204.242	-0.7442426	18585.48	24697.92
3	TESLA INC	413.256	-0.7357798	18585.48	24697.92
4	TESLA INC	2013.496	-0.6709872	18585.48	24697.92
5	TESLA INC	3198.356	-0.6230132	18585.48	24697.92
6	TESLA INC	4046.025	-0.5886917	18585.48	24697.92
7	TESLA INC	7000.132	-0.4690821	18585.48	24697.92
8	TESLA INC	11758.751	-0.2764092	18585.48	24697.92
9	TESLA INC	21461.268	0.1164384	18585.48	24697.92
10	TESLA INC	24578.000	0.2426325	18585.48	24697.92